

## Apply logistic regression for the given data set and analyse the performance of the algorithm

Data Set: <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>

### Predicting Breast Cancer - Logistic Regression

#### Loading dataset:

```
# import dependencies
# data cleaning and manipulation
import pandas as pd
import numpy as np
# data visualization
import matplotlib.pyplot as plt
import seaborn as sns
# machine learning
from sklearn.preprocessing import StandardScaler
import sklearn.linear_model as skl_lm
from sklearn import preprocessing
from sklearn import neighbors
from sklearn.metrics import confusion_matrix, classification_report, precision_score
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
import statsmodels.formula.api as smf

# initialize some package settings
sns.set(style="whitegrid", color_codes=True, font_scale=1.3)
#%matplotlib inline
# read in the data and check the first 5 rows
df = pd.read_csv('./breastcancer.csv', index_col=0)
print(df.head())

# general summary of the dataframe
print(df.info())
```

id			...		
842302	M	17.99	...	0.11890	NaN
842517	M	20.57	...	0.08902	NaN
84300903	M	19.69	...	0.08758	NaN
84348301	M	11.42	...	0.17300	NaN
84358402	M	20.29	...	0.07678	NaN

```

Data columns (total 32 columns):
diagnosis                569 non-null object
radius_mean              569 non-null float64
texture_mean             569 non-null float64
perimeter_mean           569 non-null float64
area_mean                569 non-null float64
smoothness_mean          569 non-null float64
compactness_mean         569 non-null float64
concavity_mean           569 non-null float64
concave points_mean      569 non-null float64
symmetry_mean            569 non-null float64
fractal_dimension_mean   569 non-null float64
radius_se                569 non-null float64
texture_se               569 non-null float64
perimeter_se             569 non-null float64
area_se                  569 non-null float64
smoothness_se            569 non-null float64
compactness_se           569 non-null float64
concavity_se             569 non-null float64
concave points_se        569 non-null float64
symmetry_se              569 non-null float64
fractal_dimension_se     569 non-null float64
radius_worst             569 non-null float64
texture_worst            569 non-null float64
perimeter_worst          569 non-null float64
area_worst               569 non-null float64
smoothness_worst         569 non-null float64
compactness_worst        569 non-null float64
concavity_worst          569 non-null float64
concave points_worst     569 non-null float64
symmetry_worst           569 non-null float64
fractal_dimension_worst  569 non-null float64
Unnamed: 32              0 non-null float64
dtypes: float64(31), object(1)
memory usage: 146.7+ KB
None

```

## Removing missing value column:

```

# remove the 'Unnamed: 32' column
df = df.drop('Unnamed: 32', axis=1)
# check the data type of each column
print(df.dtypes)

```

```

diagnosis                object
radius_mean              float64
texture_mean             float64
perimeter_mean           float64
area_mean                float64
smoothness_mean          float64
compactness_mean         float64
concavity_mean           float64
concave points_mean      float64
symmetry_mean            float64
fractal_dimension_mean   float64
radius_se                float64
texture_se               float64
perimeter_se             float64
area_se                  float64
smoothness_se            float64
compactness_se           float64
concavity_se             float64
concave points_se        float64
symmetry_se              float64
fractal_dimension_se     float64
radius_worst             float64
texture_worst            float64
perimeter_worst          float64
area_worst               float64
smoothness_worst         float64
compactness_worst        float64
concavity_worst          float64
concave points_worst     float64
symmetry_worst           float64
fractal_dimension_worst  float64


```

## Distribution of response variable classes:

```
# visualize distribution of classes
plt.figure(figsize=(8, 4))
sns.countplot(df['diagnosis'], palette='RdBu')
# count number of obs in each class
benign, malignant = df['diagnosis'].value_counts()
print('Number of cells labeled Benign: ', benign)
print('Number of cells labeled Malignant : ', malignant)
print('')
print('% of cells labeled Benign', round(benign / len(df) * 100, 2), '%')
print('% of cells labeled Malignant', round(malignant / len(df) * 100, 2), '%')
```

```
Number of cells labeled Benign: 357
Number of cells labeled Malignant : 212

% of cells labeled Benign 62.74 %
% of cells labeled Malignant 37.26 %
```



## Drop all unnecessary columns:

```
# first, drop all "worst" columns
cols = ['radius_worst',
        'texture_worst',
        'perimeter_worst',
        'area_worst',
        'smoothness_worst',
        'compactness_worst',
        'concavity_worst',
        'concave points_worst',
        'symmetry_worst',
        'fractal_dimension_worst']
df = df.drop(cols, axis=1)
# then, drop all columns related to the "perimeter" and "area" attributes
cols = ['perimeter_mean',
        'perimeter_se',
        'area_mean',
        'area_se']
df = df.drop(cols, axis=1)
# lastly, drop all columns related to the "concavity" and "concave points" attributes
cols = ['concavity_mean',
        'concavity_se',
        'concave points_mean',
        'concave points_se']
df = df.drop(cols, axis=1)
# verify remaining columns
print(df.columns)
```

```
Index(['diagnosis', 'radius_mean', 'texture_mean', 'smoothness_mean',
      'compactness_mean', 'symmetry_mean', 'fractal_dimension_mean',
      'radius_se', 'texture_se', 'smoothness_se', 'compactness_se',
      'symmetry_se', 'fractal_dimension_se'],
      dtype='object')
```

### Splitting data set to reduce overfitting:

```
# Split the data into training and testing sets
X = df
y = df['diagnosis']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=40)
```

**Formula to be used for the logistic regression:**

```
# Create a string for the formula
cols = df.columns.drop('diagnosis')
formula = 'diagnosis ~ ' + ' + '.join(cols)
print(formula, '\n')
```

```
diagnosis ~ radius_mean + texture_mean + smoothness_mean + compactness_mean + symmetry_mean +
fractal_dimension_mean + radius_se + texture_se + smoothness_se + compactness_se + symmetry_se +
fractal_dimension_se
```

```
# Run the model and report the results
model = smf.glm(formula=formula, data=X_train, family=sm.families.Binomial())
logistic_fit = model.fit()
print(logistic_fit.summary())
```

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	['diagnosis[B]', 'diagnosis[M]']		No. Observations:	398		
Model:	GLM		Df Residuals:	385		
Model Family:	Binomial		Df Model:	12		
Link Function:	logit		Scale:	1.0000		
Method:	IRLS		Log-Likelihood:	-55.340		
Date:	Sun, 02 Feb 2020		Deviance:	110.68		
Time:	17:06:33		Pearson chi2:	125.		
No. Iterations:	9					
Covariance Type:	nonrobust					
=====						
	coef	std err	z	P> z	[0.025	0.975]
-----						
Intercept	44.5427	11.787	3.779	0.000	21.441	67.644
radius_mean	-1.1610	0.301	-3.862	0.000	-1.750	-0.572
texture_mean	-0.4237	0.087	-4.866	0.000	-0.594	-0.253
smoothness_mean	-85.3981	40.976	-2.084	0.037	-165.709	-5.088
compactness_mean	-16.7104	22.510	-0.742	0.458	-60.829	27.408
symmetry_mean	-46.2721	17.767	-2.604	0.009	-81.095	-11.449
fractal_dimension_mean	-49.1536	121.888	-0.403	0.687	-288.050	189.742
radius_se	-7.1916	2.806	-2.563	0.010	-12.691	-1.692
texture_se	0.1849	0.784	0.236	0.814	-1.353	1.722
smoothness_se	163.6068	159.702	1.024	0.306	-149.403	476.616
compactness_se	-31.1808	42.772	-0.729	0.466	-115.012	52.650
symmetry_se	74.7366	51.458	1.452	0.146	-26.119	175.592
fractal_dimension_se	824.1245	412.040	2.000	0.045	16.541	1631.708
=====						

## Prediction:

```
# predict the test data and show the first 5 predictions
predictions = logistic_fit.predict(X_test)
print(predictions[1:6])
```

```
id
848406      0.324251
907915      0.996906
911201      0.964710
84799002     0.000544
8911164     0.838719
dtype: float64
```

```
# Note how the values are numerical.
# Convert these probabilities into nominal values and check the first 5 predictions
again.
predictions_nominal = [ "M" if x < 0.5 else "B" for x in predictions]
print(predictions_nominal[1:6])
```

```
['M', 'B', 'B', 'M', 'B']
```

We can confirm that probabilities closer to 0 have been labeled as "M", while the ones closer to 1 have been labelled as "B". Now we are able to evaluate the accuracy of our predictions by checking out the classification report and the confusion matrix.

```
print(classification_report(y_test, predictions_nominal, digits=3))
cfm = confusion_matrix(y_test, predictions_nominal)

true_negative = cfm[0][0]
false_positive = cfm[0][1]
false_negative = cfm[1][0]
true_positive = cfm[1][1]

print('Confusion Matrix: \n', cfm, '\n')

print('True Negative:', true_negative)
print('False Positive:', false_positive)
print('False Negative:', false_negative)
print('True Positive:', true_positive)
print('Correct Predictions',
      round((true_negative + true_positive) / len(predictions_nominal) * 100, 1), '%')
```

	precision	recall	f1-score	support
B	0.982	0.965	0.974	115
M	0.931	0.964	0.947	56
accuracy			0.965	171
macro avg	0.957	0.965	0.961	171
weighted avg	0.966	0.965	0.965	171

```
Confusion Matrix:
[[111  4]
 [ 2 54]]
```

```
True Negative: 111
False Positive: 4
False Negative: 2
True Positive: 54
Correct Predictions 96.5 %
```

The model has accurately labelled 96.5% of the test data